

Amazon Athena

Michael Lin
Senior Solutions Architect
Amazon Web Services

Agenda

- Introduction to Amazon Athena
- Athena Design Patterns
- Athena in Action
- Summary

1. Introduction to Amazon Athena

Amazon Athena

Amazon Athena is an interactive query service that makes it easy to analyze data using standard SQL. Athena is serverless, so there is no infrastructure to manage, and you pay only for the queries that you run.

- Query data in your Amazon S3 based data lake
- Analyze infrastructure, operation, and application logs
- Interactive analytics using popular BI tools
- Self-service data exploration for data scientists
- Embed analytics capabilities into your applications
- Perform ETL using SQL

Why Amazon Athena

- Decouple storage from compute
- Serverless – No infrastructure or resources to manage
- Pay only for data scanned
- Schema on read – Same data, many views
- Secure – IAM/SAMLv2 for authentication; Encryption at rest & in transit
- Standard compliant and open storage file formats
- Built on powerful community supported OSS solutions

Simple Pricing

- DDL operations – FREE
- SQL operations – FREE
- Query concurrency – FREE
- Data scanned - \$5 / TB
- Standard S3 rates for storage, requests, and data transfers apply

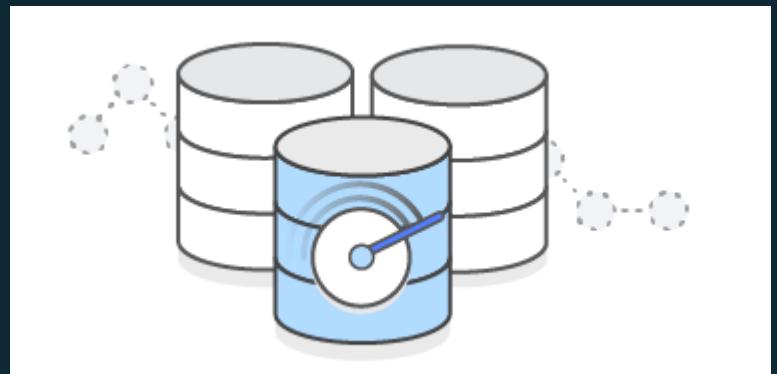
Run Standard SQL



- Uses Presto with ANSI SQL support
- Works with standard data formats
 - CSV
 - Apache Weblogs
 - JSON
 - Parquet
 - ORC
- Handles complex queries
 - Large Joins
 - Window functions
 - Arrays

Fast Performance for Large Data Sets

- Fast, ad-hoc queries
- Executes queries in parallel
- No provisioning extra resources for complex queries
- Scales automatically
- Amazon Athena Federated Query



2. Athena Design Patterns

When to use Athena vs EMR

Athena

- Ad-hoc querying
- Serverless - no setup or management of cluster
- Run queries using standard SQL
- Scales automatically based on complexity of queries

EMR

- Data processing/ETL
- Build and manage your own cluster
- Run custom applications and code
- Use big data processing frameworks (Spark, Hadoop, Presto, or Hbase)

New — Amazon Athena for Apache Spark

by Donnie Prakoso | on 30 NOV 2022 | in [Amazon Athena](#), [Analytics](#), [Announcements](#), [AWS Re:Invent](#), [Launch](#), [News](#) |

[Permalink](#) | [Comments](#) | [Share](#)

▶ 0:00 / 0:00



Voiced by [Amazon Polly](#)

When Jeff Barr first [announced Amazon Athena](#) in 2016, it changed my perspective on interacting with data. With [Amazon Athena](#), I can interact with my data in just a few steps—starting from creating a table in Athena, loading data using connectors, and querying using the ANSI SQL standard.

Over time, various industries, such as financial services, healthcare, and retail, have needed to run more complex analyses for a variety of formats and sizes of data. To facilitate complex data analysis, organizations adopted [Apache Spark](#). Apache Spark is a popular, open-source, distributed processing system designed to run fast analytics workloads for data of any size.

<https://aws.amazon.com/blogs/aws/new-amazon-athena-for-apache-spark/>

When to use Athena vs Redshift

Athena

- Ad-hoc querying
- Serverless - no setup or management of cluster
- Run queries using standard SQL
- Scales automatically based on complexity of queries

Redshift

- Data warehouse (historical analysis and reporting)
- Need to setup a cluster
- Run queries against highly structured data with many joins
- Can use same S3 data source as Athena

Amazon Redshift Spectrum Extends Data Warehousing Out to Exabytes—No Loading Required

by Maor Kleider | on 21 JUL 2017 | in [Amazon Redshift](#), [AWS Big Data](#) | [Permalink](#) | [Comments](#) | [Share](#)

When we first looked into the possibility of building a cloud-based data warehouse many years ago, we were struck by the fact that our customers were storing ever-increasing amounts of data, and yet only a small fraction of that data ever made it into a data warehouse or Hadoop system for analysis. We saw that this wasn't just a cloud-specific anomaly. It was also true in the broader industry, where the growth rate of the enterprise storage market segment greatly surpassed that of the data warehousing market segment.

<https://aws.amazon.com/blogs/big-data/amazon-redshift-spectrum-extends-data-warehousing-out-to-exabytes-no-loading-required/>

Configure and optimize performance of Amazon Athena federation with Amazon Redshift

by Harsha Tadiparthi | on 26 AUG 2020 | in [Amazon Athena](#), [Amazon Redshift](#), [Amazon Redshift](#), [Analytics](#), [AWS Big Data](#), [AWS Lambda](#), [Database](#) | [Permalink](#) | [Comments](#) | [Share](#)

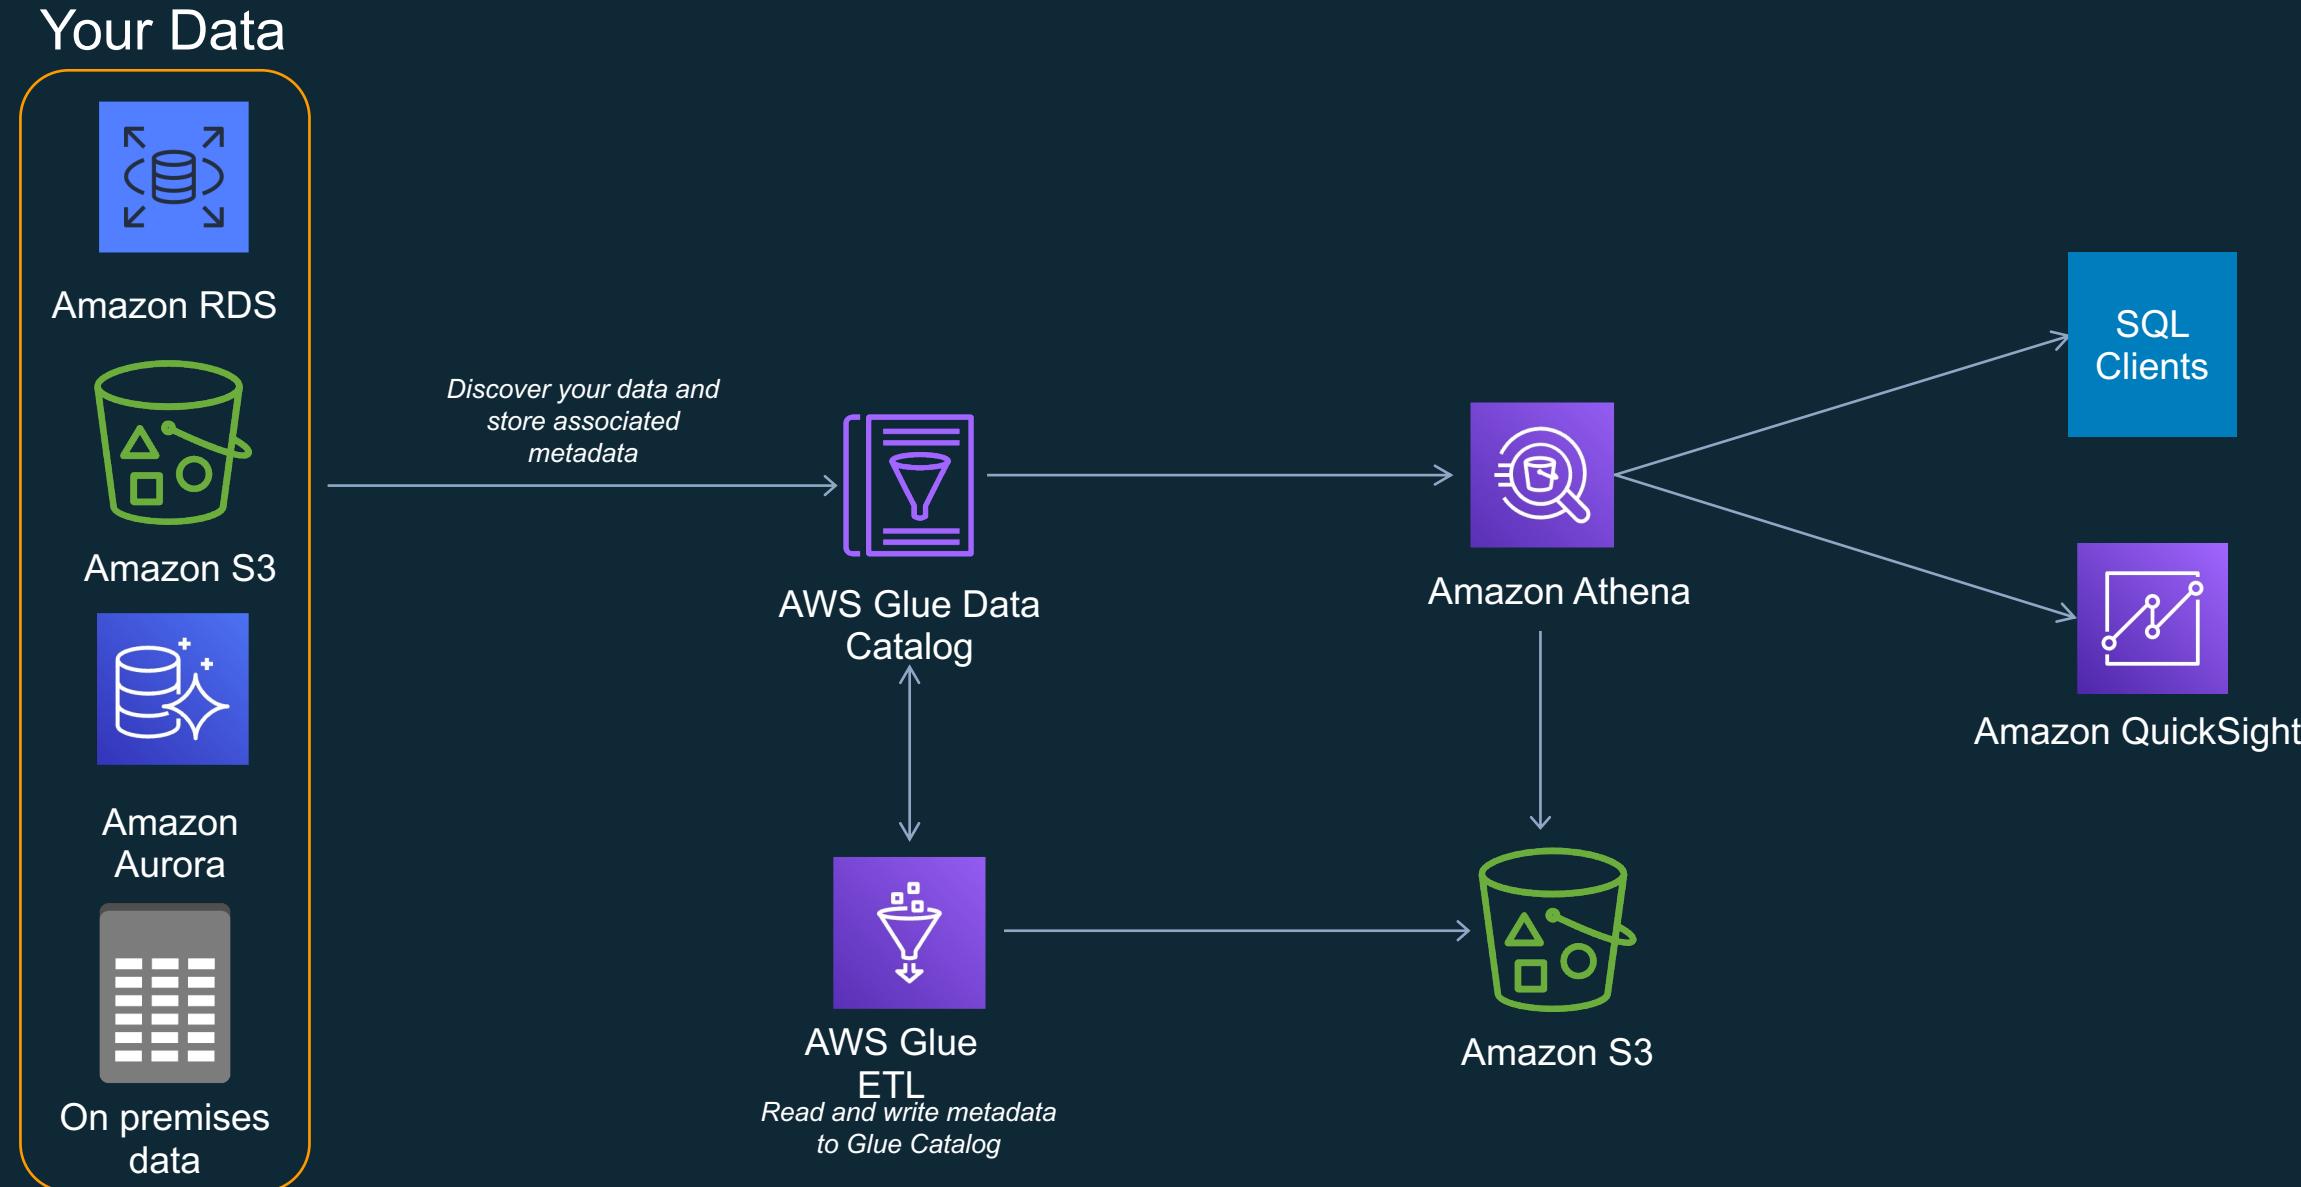
This post provides guidance on how to configure [Amazon Athena](#) federation with [AWS Lambda](#) and [Amazon Redshift](#), while addressing performance considerations to ensure proper use.

If you use data lakes in [Amazon Simple Storage Service](#) (Amazon S3) and use Amazon Redshift as your data warehouse, you may want to integrate the two for a lake house approach. Lake House is the ability to integrate Data Lake and Data warehouse seamlessly. When you need to query your data lake from your Amazon Redshift Data warehouse, you can use [Amazon Redshift Spectrum](#), which works great in unifying your data lake and data warehouse. However, when you use Athena in the data lake and need to access data in Amazon Redshift for the following two scenarios which are commonly seen, there is no easy approach:

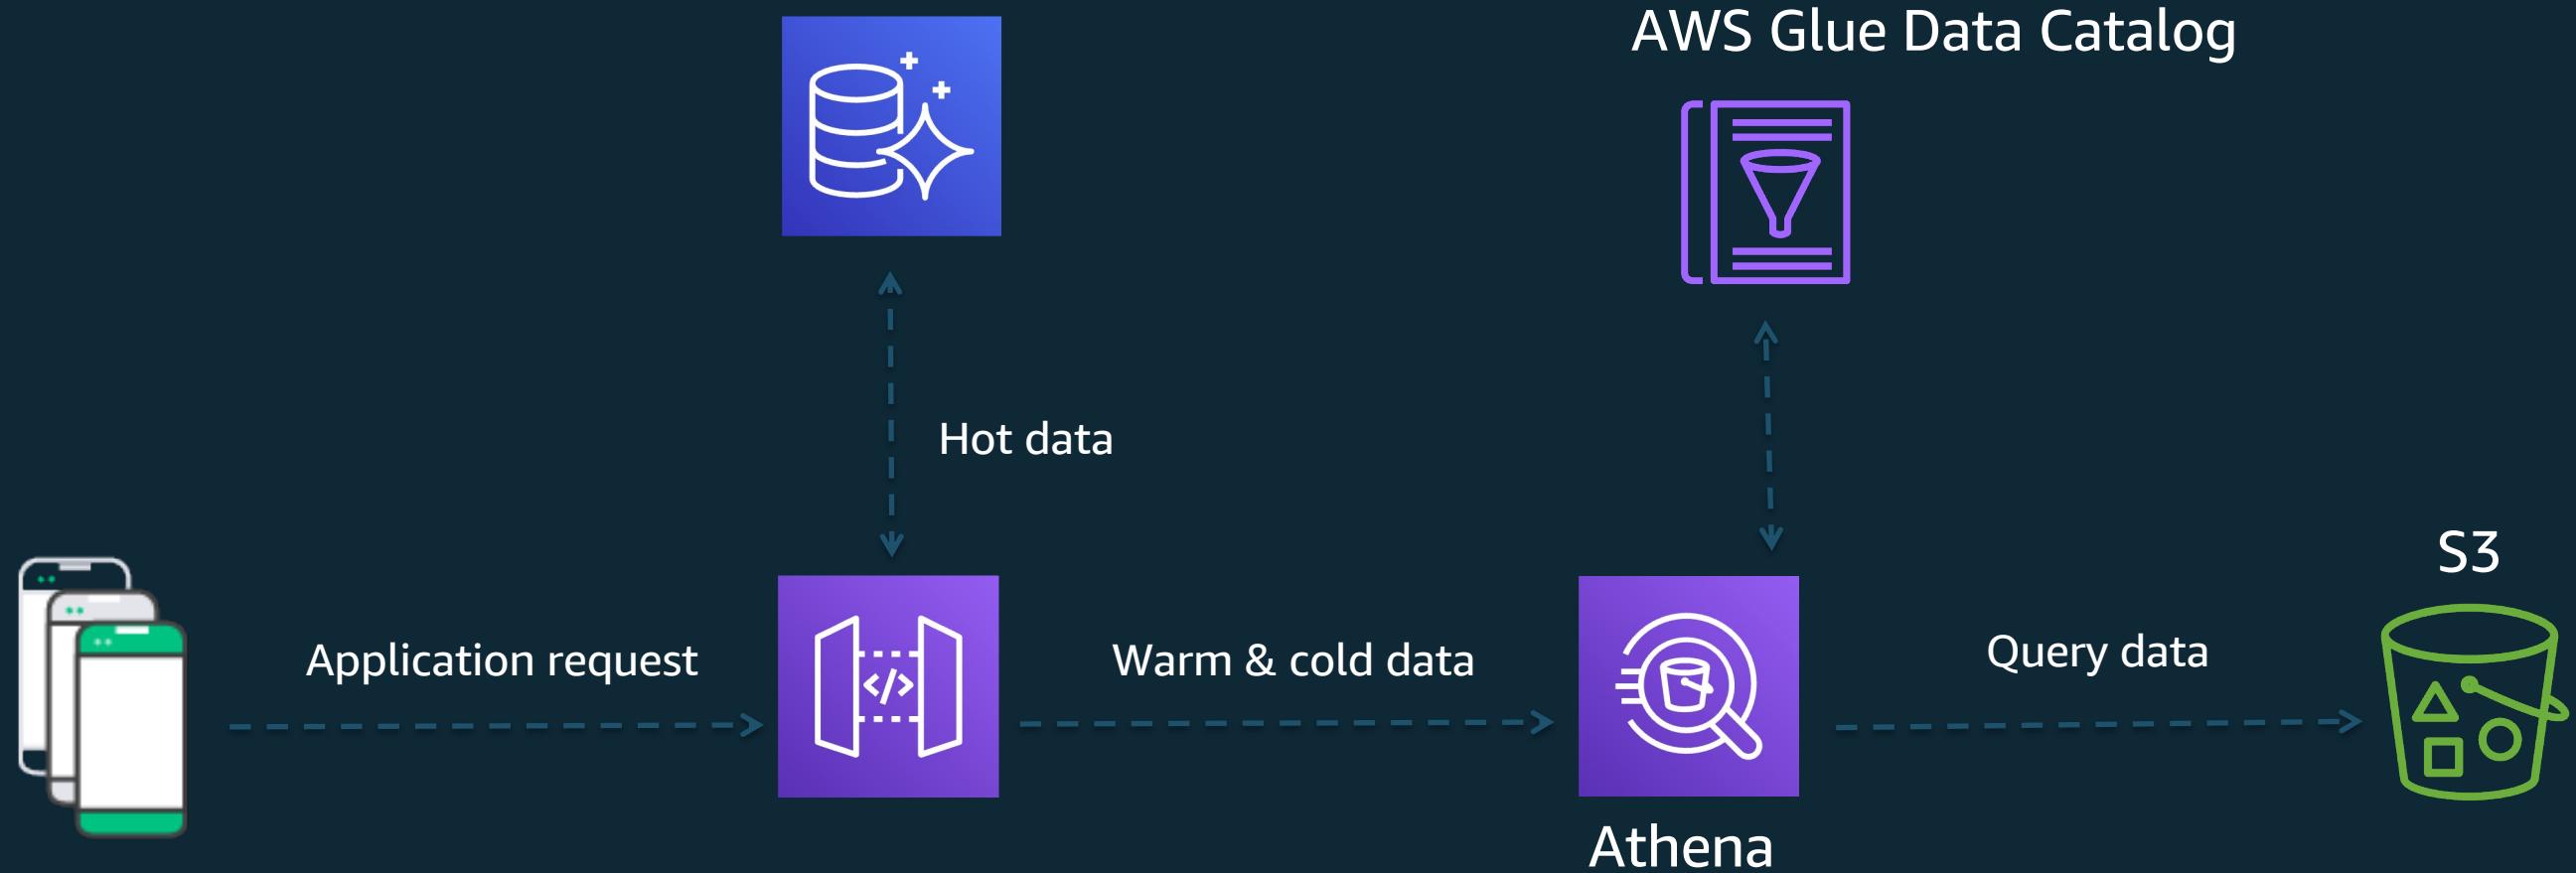
- Team A has a data lake in Amazon S3 and uses Athena. They need access to the data in an Amazon Redshift cluster owned by Team B.
- Analysts using Athena to query their data lake for analytics need agility and flexibility to access data in an Amazon Redshift data warehouse without moving the data to Amazon S3 Data Lake.

<https://aws.amazon.com/blogs/big-data/configure-and-optimize-performance-of-amazon-athena-federation-with-amazon-redshift/>

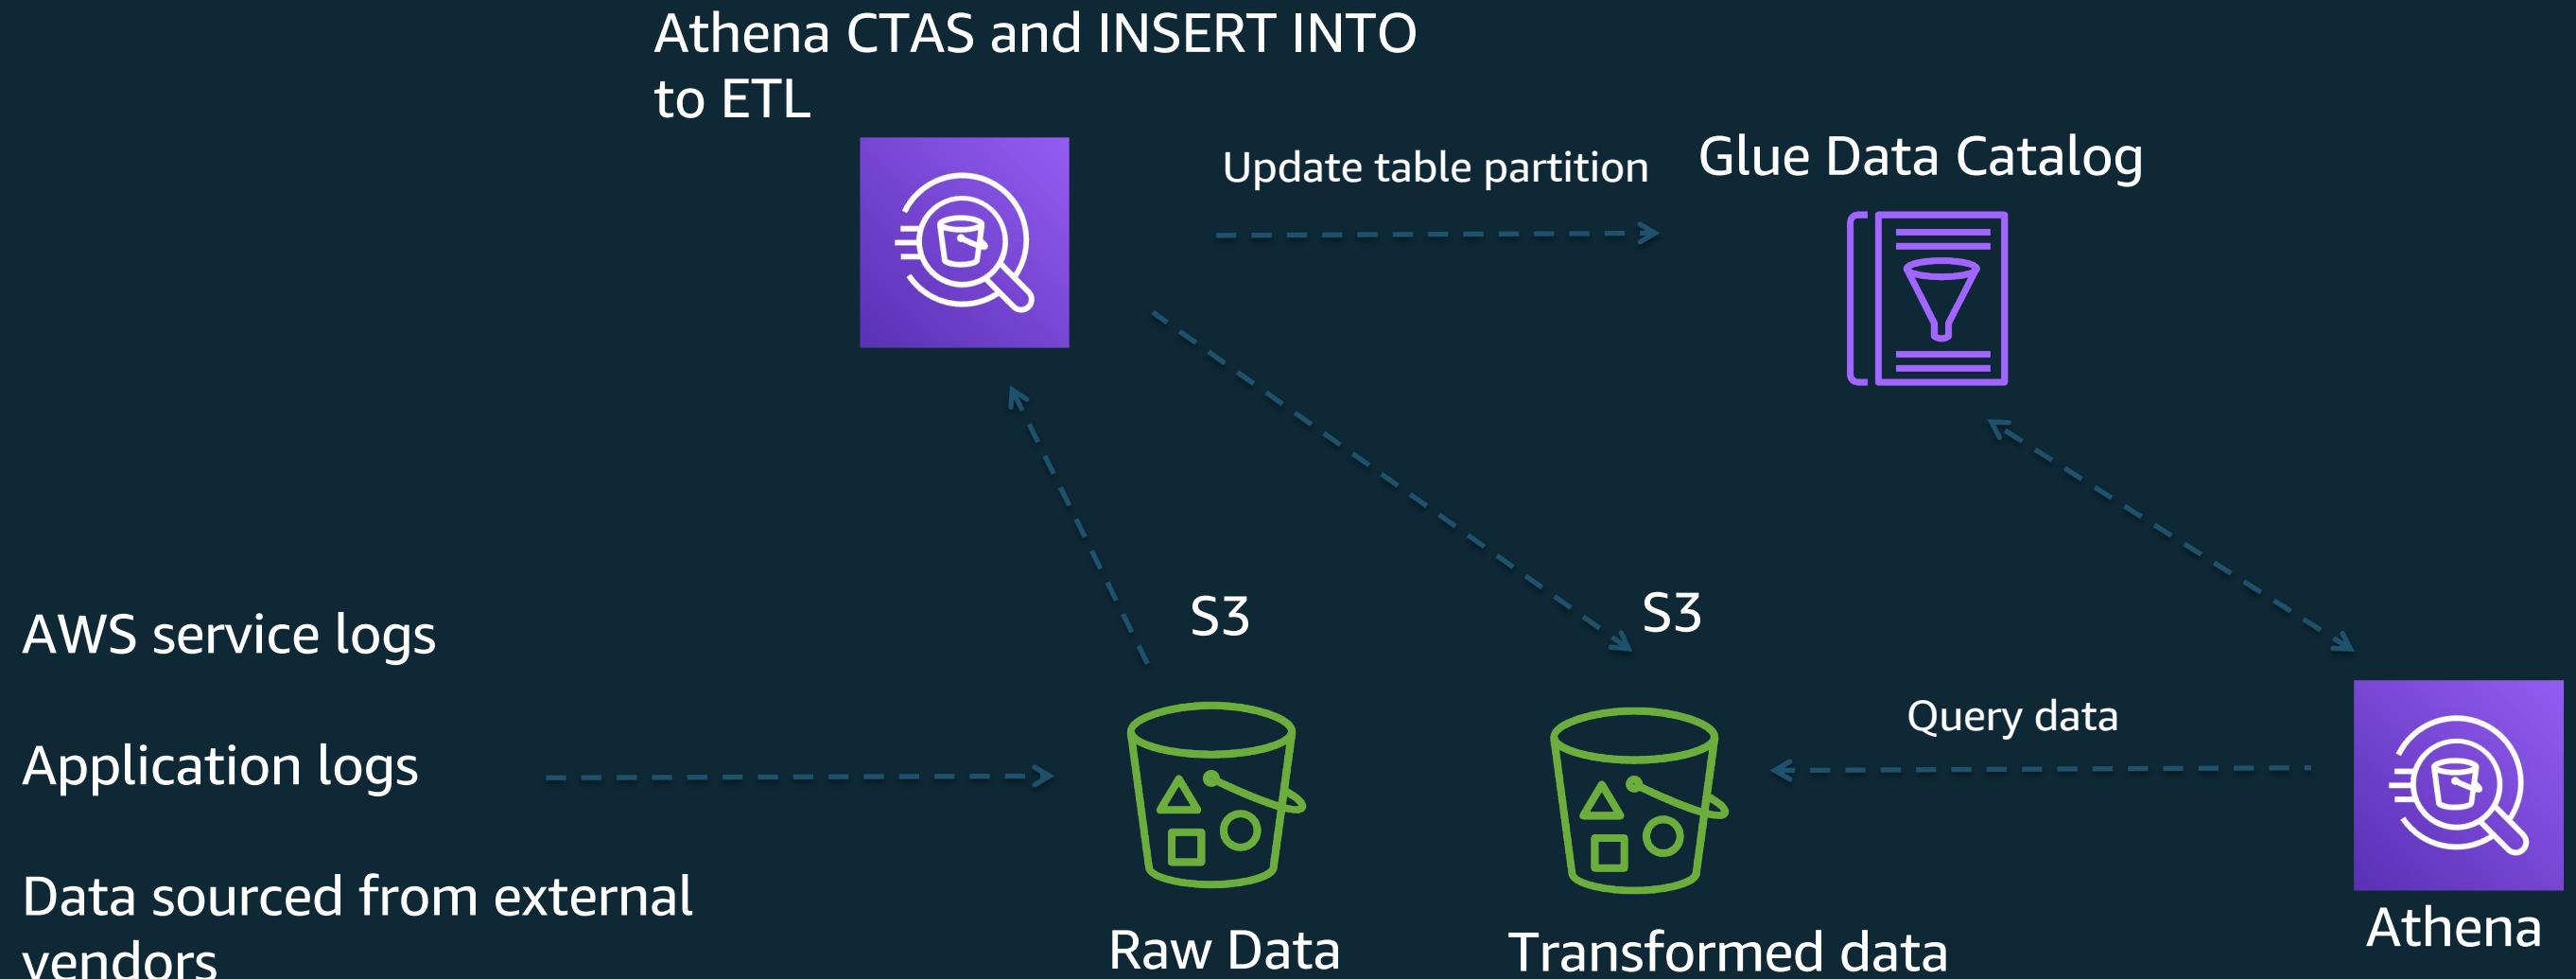
1: Data Exploration use-case



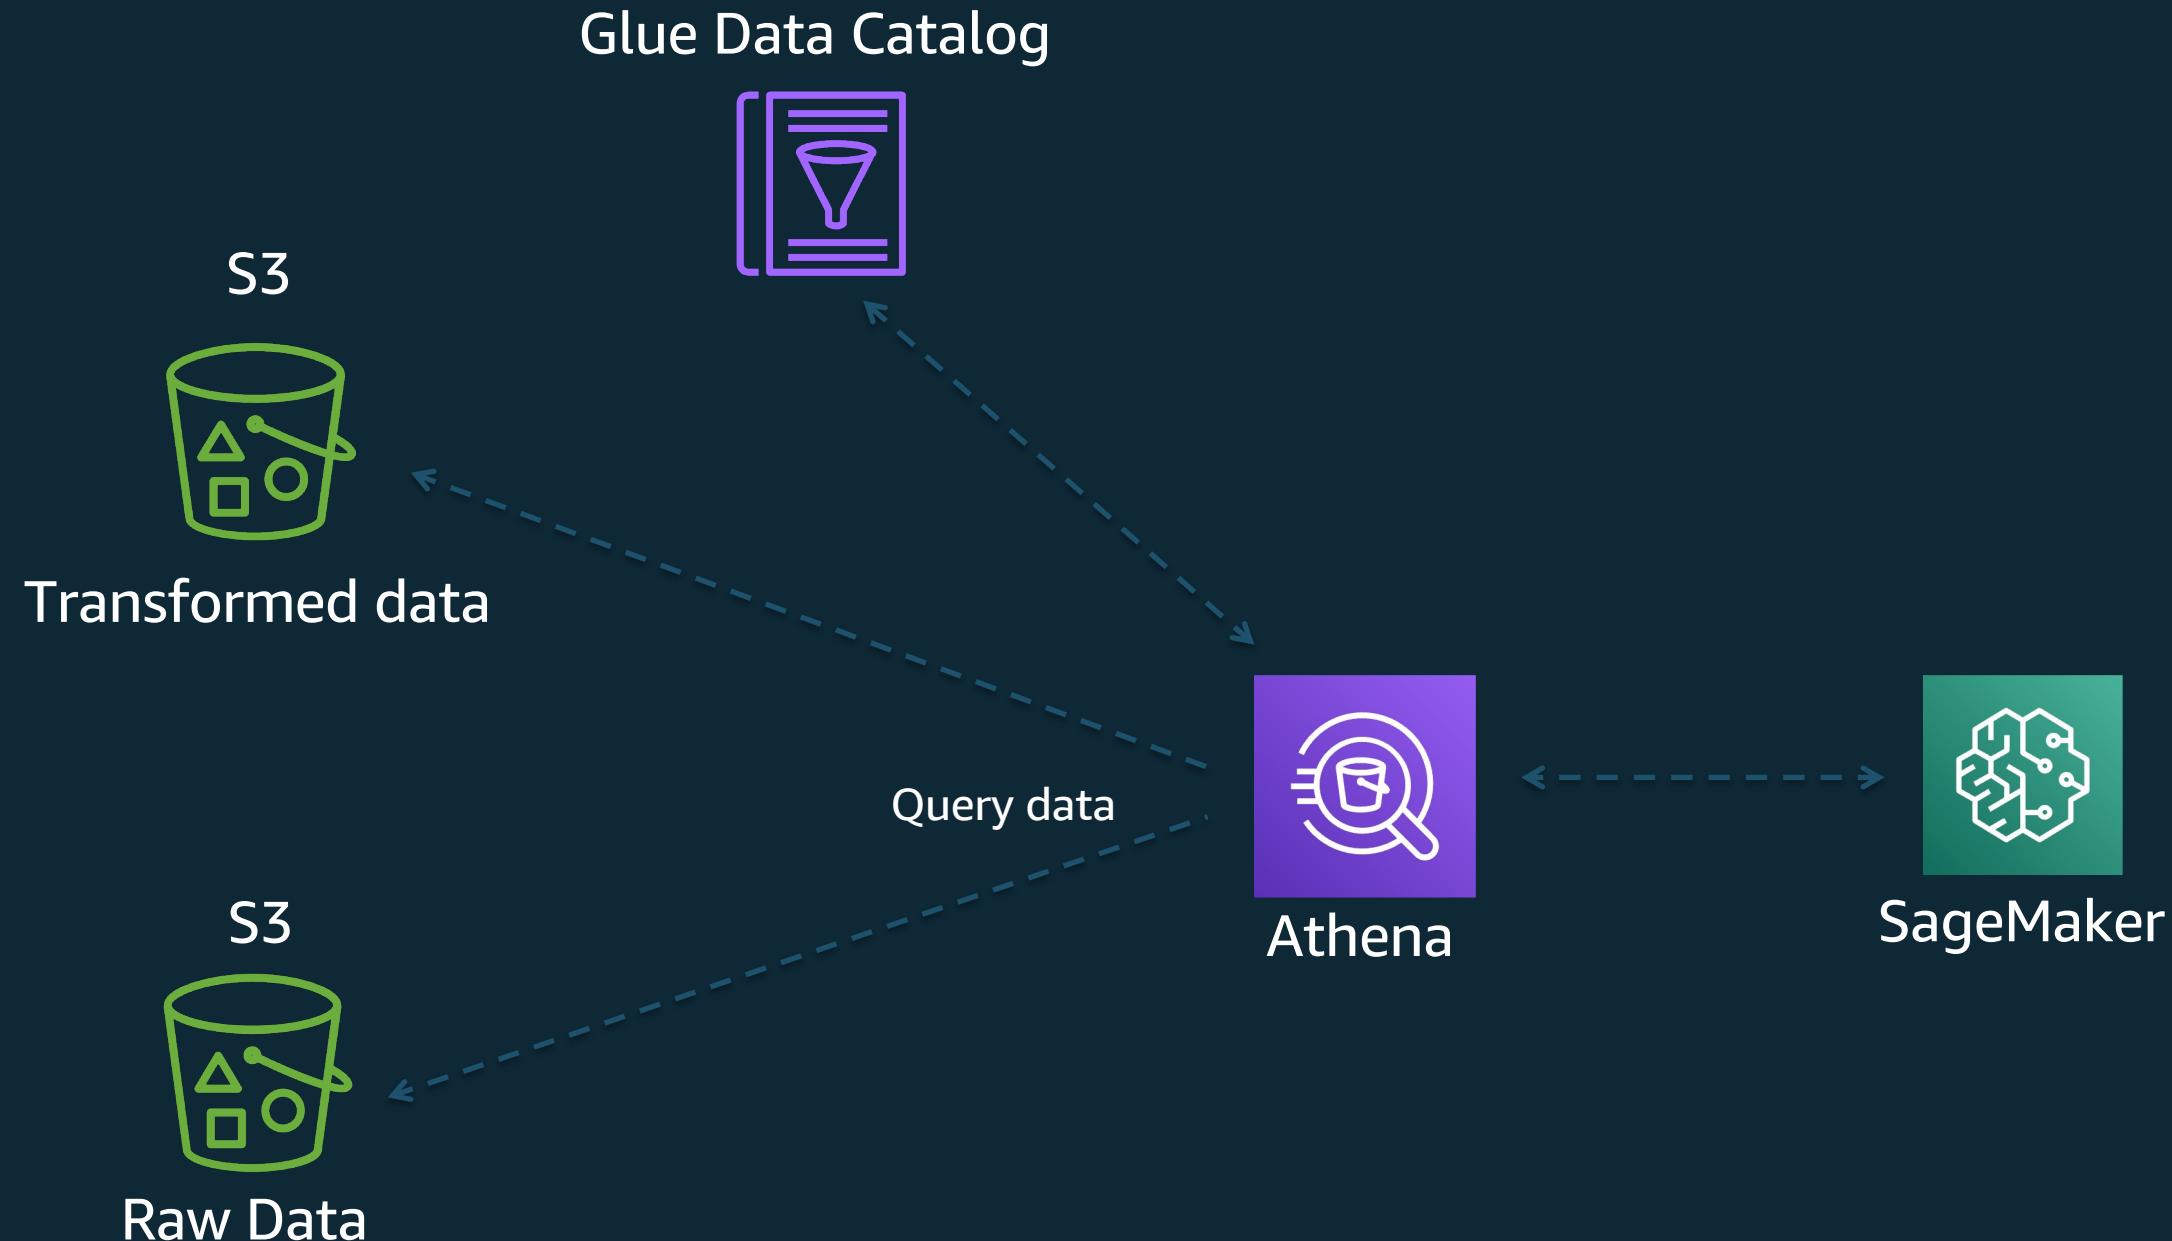
2: SaaS use-case



3: ETL and query use-case



4: Data science exploration and feature engineering



Querying AWS service logs

Topics

- [Application Load Balancer](#)
- [Elastic Load Balancing](#)
- [CloudFront](#)
- [CloudTrail](#)
- [Amazon EMR](#)
- [Global Accelerator](#)
- [GuardDuty](#)
- [Network Firewall](#)
- [Network Load Balancer](#)
- [Route 53](#)
- [Amazon SES](#)
- [Amazon VPC](#)
- [AWS WAF](#)

<https://docs.aws.amazon.com/athena/latest/ug/querying-aws-service-logs.html>

3. Athena in Action

Create External Tables

- Use Apache Hive DDL to create table
 - Run DDL statements using the Athena console
 - Via a JDBC driver, using SQL workbench
 - Using the Athena create table wizard
- Create “external” table in DDL
 - Creates a view of the data
 - Deleting table doesn’t delete the data in S3
- Schema-on-read
 - Projects your schema onto data at query execution time
 - No need for data loading or transformation

Create External Table using DDL

The screenshot shows the AWS Athena Query Editor interface. On the left, the sidebar displays the selected database ('sampledb') and lists three tables: 'elb_logs', 'flights_parquet (Partitioned)', and 'test_table (Partitioned)'. Below the tables, it indicates 'Views (0)' and provides instructions to create a view. The main area contains a query editor window titled 'New query 1' with the following DDL code:

```
1 CREATE EXTERNAL TABLE IF NOT EXISTS flights_parquet (
2     yr INT,
3     quarter INT,
4     month INT,
5     dayofmonth INT,
6     dayofweek INT,
7     flightdate STRING,
8     uniquecarrier STRING,
9     airlineid INT,
10    carrier STRING,
11    tailnum STRING,
12    flightnum STRING,
13    originairportid INT,
14    originairportseqid INT,
15    origincitymarketid INT,
16    origin STRING,
17    origincityname STRING,
18    originstate STRING,
19    originstatefips STRING,
20    originstatename STRING,
```

Below the code, there are several action buttons: 'Run query' (highlighted in blue), 'Save as', 'Create', 'Format query', and 'Clear'. A note at the bottom says 'Use Ctrl + Enter to run query, Ctrl + Space to autocomplete'. The bottom section is labeled 'Results'.

New Flights Parquet Table Created

Athena **Query Editor** Saved Queries History Catalog Manager Settings Tutorial Help

DATABASE sampledb TABLES Filter Tables... Add table...

```
1 SELECT * FROM flights_parquet limit 100;
```

Run Query Save As Format Query New Query (Run time: 8.62 seconds, Data scanned: 7.7MB) Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

yr	quarter	month	dayofmonth	dayofweek	flightdate	uniquecarrier	airlineid	carrier	tailnum	flightnum	originairportid	originairportseqid	origincitymarketid	origin	origincityname	originstate	originstatefips	originstatename	
1	2010	4	10	14	4	2010-10-14	"AA"	19805	"AA"	"N3FDAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
2	2010	4	10	15	5	2010-10-15	"AA"	19805	"AA"	"N3GEAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
3	2010	4	10	16	6	2010-10-16	"AA"	19805	"AA"	"N3GEAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
4	2010	4	10	17	7	2010-10-17	"AA"	19805	"AA"	"N3GAAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
5	2010	4	10	18	1	2010-10-18	"AA"	19805	"AA"	"N3GAAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
6	2010	4	10	19	2	2010-10-19	"AA"	19805	"AA"	"N3DNAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
7	2010	4	10	20	3	2010-10-20	"AA"	19805	"AA"	"N3GBAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
8	2010	4	10	21	4	2010-10-21	"AA"	19805	"AA"	"N3DVAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
9	2010	4	10	22	5	2010-10-22	"AA"	19805	"AA"	"N3DVAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
10	2010	4	10	23	6	2010-10-23	"AA"	19805	"AA"	"N3EHAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
11	2010	4	10	24	7	2010-10-24	"AA"	19805	"AA"	"N3EKAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
12	2010	4	10	25	1	2010-10-25	"AA"	19805	"AA"	"N3FXAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
13	2010	4	10	26	2	2010-10-26	"AA"	19805	"AA"	"N3FJAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
14	2010	4	10	27	3	2010-10-27	"AA"	19805	"AA"	"N3EHAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
15	2010	4	10	28	4	2010-10-28	"AA"	19805	"AA"	"N3EHAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
16	2010	4	10	29	5	2010-10-29	"AA"	19805	"AA"	"N3FEAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"
17	2010	4	10	30	6	2010-10-30	"AA"	19805	"AA"	"N3FEAA"	"1605"	13930	1393001	30977	"ORD"	"Chicago, IL"	"IL"	"17"	"Illinois"

Navigate to 'Saved Queries' to get DDL

The screenshot shows the AWS Athena interface with the 'Saved Queries' tab selected. The page displays a list of saved queries, each with a name, description, and the corresponding Hive DDL or MSCK repair command.

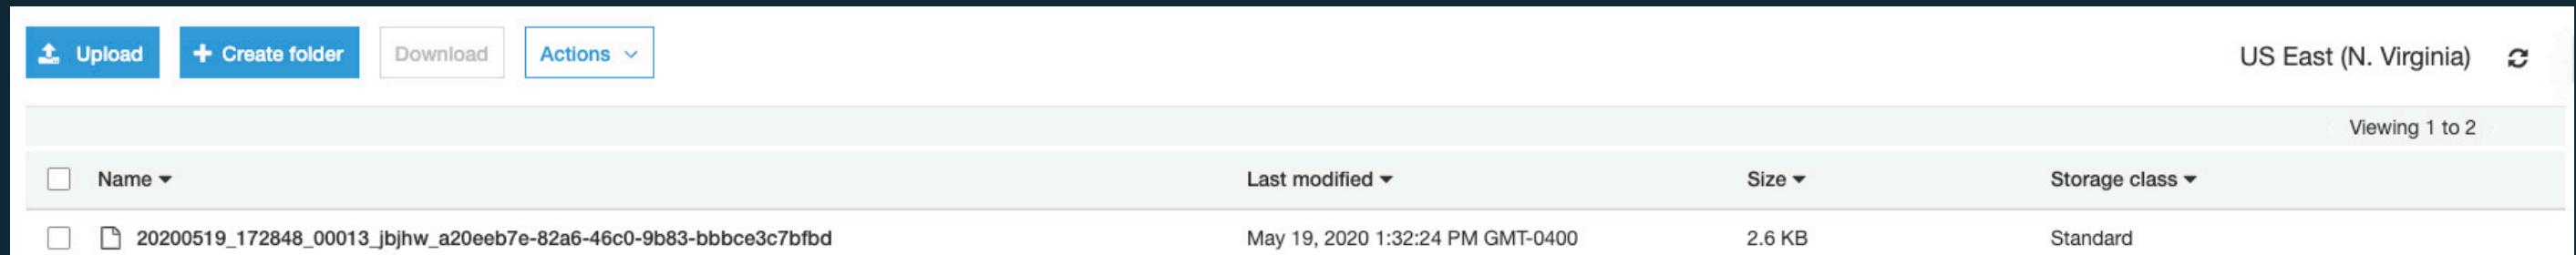
Name	Description	Query
Load flights table partitions	Sample query to load flights table partitions using MSCK REP...	MSCK REPAIR TABLE flights_parquet;
ELB Select Query	Sample query to view peak load ELBs during a particular time...	SELECT elb_name, count(1) FROM elb_logs_orc Where elb_resp...
Create CloudFront Table	Sample Hive DDL statement to create a table pointing to Clou...	CREATE EXTERNAL TABLE IF NOT EXISTS cloudfront_logs (Date...
Create table with partitions	Sample Hive DDL statement to create a partitioned table poin...	CREATE EXTERNAL TABLE IF NOT EXISTS flights_parquet (yr I...
CloudFront Select Query	Sample query to view requests per operating system during a ...	SELECT os, COUNT(*) count FROM cloudfront_logs WHERE date BE...
Create ELB Table	Sample Hive DDL statement to create a table pointing to ELB ...	CREATE EXTERNAL TABLE IF NOT EXISTS elb_logs_orc (request...
Flights Select Query	Sample query to get the top 10 airports with the most number...	SELECT origin, count(*) AS total_departures FROM flights_par...

Save Query Results

- When you run an Athena query for the first time, an S3 bucket called "aws-athena-query-results-<account_id>" is created on your account, where "<account_id>" is replaced with your AWS account ID.

```
aws-athena-query-results-{account_id}-{region}/Unsaved/{year}/{month}/{day}/{query_id}.csv
```

- The results of all Athena queries are stored in this S3 bucket using the year, month, and day the query was run, the hexadecimal Athena query ID, and the region the query was run in the following format:



Viewing 1 to 2			
Name	Last modified	Size	Storage class
20200519_172848_00013_jbjhw_a20eeb7e-82a6-46c0-9b83-bbbce3c7bfbd.csv	May 19, 2020 1:32:24 PM GMT-0400	2.6 KB	Standard

3.1 Tuning and Design Patterns

Data Formats and SerDes

- Current data formats and SerDes
 - Standard
 - CSV, TSV
 - JSON
 - Apache web logs
 - Custom-Delimited Files
 - Open-sourced columnar formats
 - Parquet
 - ORC

Convert to Columnar Formats-Parquet/ORC

- Convert to Parquet/ORC, use EMR
 - Create EMR cluster with Hive or Spark
 - In the step section of create cluster
 - Specify script in S3, Hive or PySpark
 - Point to the input data in text/CSV/JSON
 - Create output data in Parquet/ORC in S3
- Use Athena CTAS queries to create a new table with Parquet data from a source table in a different format.
 - `CREATE TABLE new_table WITH (format = 'Parquet', parquet_compression = 'SNAPPY') AS SELECT * FROM old_table;`

Column Data Types

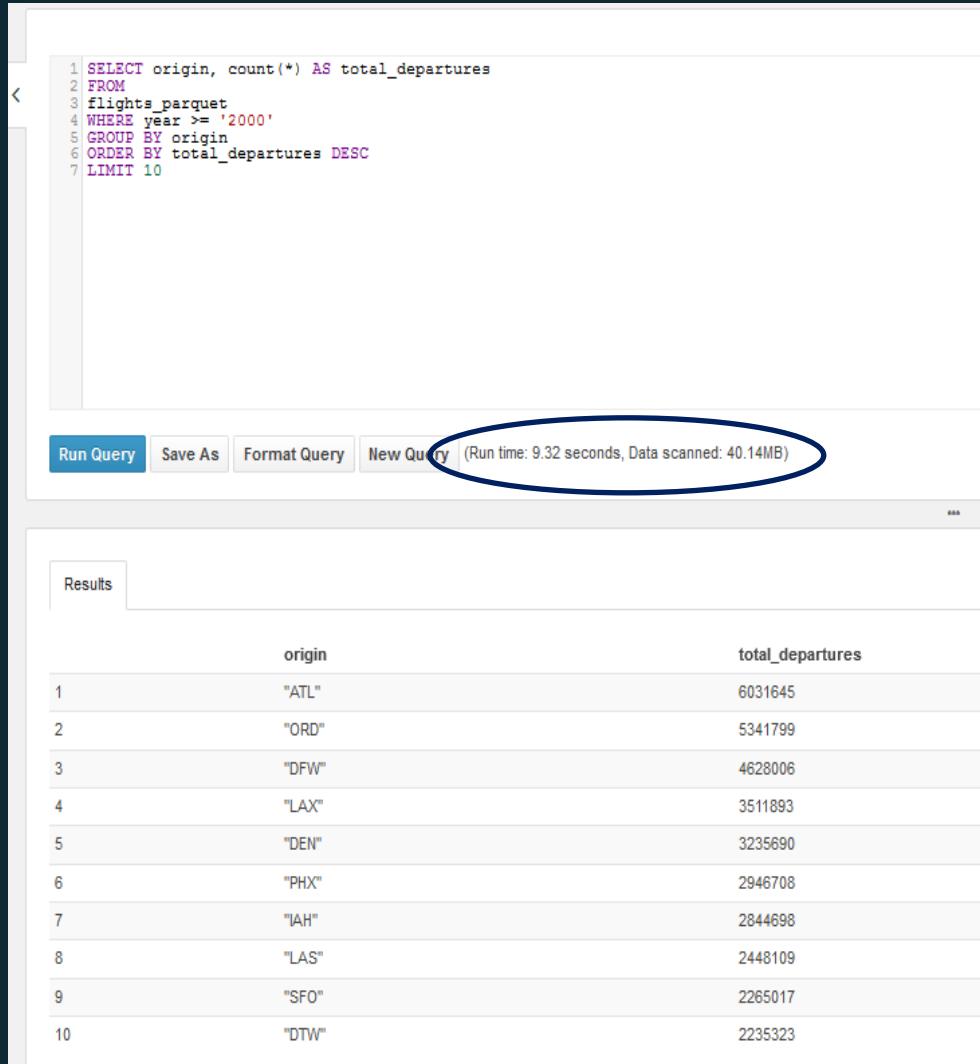
- Column Data Types
 - Primitives (Integer, Double, Date, String, Decimal, Varchar...)
 - ARRAY
 - MAP
 - STRUCT
- Nested JSON
 - Query using structures and array of structures

3.2 Partitions

Partitions

- Partition on any column in table
- Limits the amount of data each query scans
- Common practice
 - Multi-level time-based partition
- Example
 - PARTITIONED BY (year STRING, month STRING, day STRING)
 - Performance and cost savings on queries filtering by day, month and year

How does partitioning work?



```
1 SELECT origin, count(*) AS total_departures
2 FROM
3 flights_parquet
4 WHERE year >= '2000'
5 GROUP BY origin
6 ORDER BY total_departures DESC
7 LIMIT 10
```

Run Query Save As Format Query New Query (Run time: 9.32 seconds, Data scanned: 40.14MB)

origin	total_departures
1 "ATL"	6031645
2 "ORD"	5341799
3 "DFW"	4628006
4 "LAX"	3511893
5 "DEN"	3235690
6 "PHX"	2946708
7 "IAH"	2844698
8 "LAS"	2448109
9 "SFO"	2265017
10 "DTW"	2235323

- Flight dataset in Parquet
- Partitioned by year
- Group by query
 - Top 10 airports with the most departures since 2000
- Where year >= 2000
 - Athena doesn't have to scan data in partitions before year 2000
- Amount of data scanned affects pricing
 - 1 min 6s, 3.11GB for csv unpartitioned
 - 37.37 s, 2.25GB for csv partitioned

How does partitioning work?

With partitions:
CSV: 37.37 s, 2.25GB

The screenshot shows a database query interface with a code editor containing the following SQL query:

```
1 SELECT origin, count(*) AS total_departures
2 FROM
3 flights_csv
4 WHERE year >= '2000'
5 GROUP BY origin
6 ORDER BY total_departures DESC
7 LIMIT 10
```

Below the code editor, a status bar displays: "Run Query Save As Format Query New Query (Run time: 37.37 seconds, Data scanned: 2.25GB)". A blue oval highlights this status bar.

The results table has columns "origin" and "total_departures". The data is as follows:

	origin	total_departures
1	"ATL"	6031645
2	"ORD"	5341799
3	"DFW"	4628006
4	"LAX"	3511893
5	"DEN"	3235690
6	"PHX"	2946708
7	"IAH"	2844698
8	"LAS"	2448109
9	"SFO"	2265017
10	"DTW"	2235323

No partitions:
1min 6s, 3.11GB

The screenshot shows a database query interface with a code editor containing the same SQL query as the first interface, but with a different table name:

```
1 SELECT origin, count(*) AS total_departures
2 FROM
3 flights_csv_unpartitioned
4 WHERE yr >= 2000
5 GROUP BY origin
6 ORDER BY total_departures DESC
7 LIMIT 10
```

Below the code editor, a status bar displays: "Run Query Save As Format Query New Query (Run time: 1 minutes 6 seconds, Data scanned: 3.11GB)". A blue oval highlights this status bar.

The results table has columns "origin" and "total_departures". The data is identical to the first interface:

	origin	total_departures
1	"ATL"	6031645
2	"ORD"	5341799
3	"DFW"	4628006
4	"LAX"	3511893
5	"DEN"	3235690
6	"PHX"	2946708
7	"IAH"	2844698
8	"LAS"	2448109
9	"SFO"	2265017
10	"DTW"	2235323

Partition by Running MSCK Repair

- If your data is already partitioned in Hive format
 - Use MSCK repair table command.
 - To sync all new data in S3 with Hive metastore
 - Example:
 - `msck repair table flights_csv`
 - `show partitions flights_csv`

Data stored on Amazon S3 in Hive format

```
aws s3 ls s3://elasticmapreduce/samples/hive-ads/tables/impressions/  
  
PRE dt=2009-04-12-13-00/  
PRE dt=2009-04-12-13-05/  
PRE dt=2009-04-12-13-10/  
PRE dt=2009-04-12-13-15/  
PRE dt=2009-04-12-13-20/  
PRE dt=2009-04-12-14-00/  
PRE dt=2009-04-12-14-05/  
PRE dt=2009-04-12-14-10/  
PRE dt=2009-04-12-14-15/  
PRE dt=2009-04-12-14-20/  
PRE dt=2009-04-12-15-00/  
PRE dt=2009-04-12-15-05/
```

<https://docs.aws.amazon.com/athena/latest/ug/partitions.html>

Partition by Manually Adding Partitions

- If data is not partitioned in Hive format
 - Cannot use MSCK repair table command
 - Use ALTER TABLE ADD PARTITION to add each partition manually
 - Look at automating adding partitions
 - Example:
 - `ALTER TABLE elb_logs_raw_native_part ADD PARTITION (year='2015',month='01',day='01') location 's3://athena-examples/elb/raw/2015/01/01/'`

3.3 Optimize and Secure

Access Control for Athena

- Access Control
 - AWS Identity and Access Management (IAM)
 - Access Control Lists (ACLs)
 - Amazon S3 bucket policies.
- IAM policies for S3
 - IAM is natively integrated with S3
 - Grant IAM users fine-grained control to S3
 - Restrict users from querying it using Athena



Access Control Continued

- To run queries in Athena, you must have the appropriate permissions for the following:
 - Athena API actions including additional actions for Athena workgroups
 - Amazon S3 locations where the underlying data to query is stored.
 - Metadata and resources that you store in the AWS Glue Data Catalog, such as databases and tables, including additional actions for encrypted metadata.

Best Practices

- Partition your data
 - Divides your table into parts and keeps the related data together based on column values such as date, country, region, etc
- Bucket your data
 - You can specify one or more columns containing rows that you want to group together and put those rows into multiple buckets.
- Compress and split files
 - We recommend using either Apache Parquet or Apache ORC, which compress data by default and are splitable
- Optimize File Sizes
 - Ideally not less than 128 MB

Best Practices Continued

- Optimize columnar data store generation
 - Apache Parquet and Apache ORC
- Optimize Queries
 - ORDER BY (with limit)
 - JOIN (with larger table on the left)
 - GROUP BY (with higher to lower cardinality)
 - LIKE (better substituted using regex)
- Only include the columns that you need

Top 10 Performance Tuning Tips for Amazon Athena

by Mert Hocanin and Pathik Shah | on 24 MAR 2017 | in [Amazon Athena](#), [AWS Big Data](#) | [Permalink](#) | [!\[\]\(8268d3ddd38f4b31328094274c2a1f20_img.jpg\) Comments](#) |

[!\[\]\(081c33931122b4c6b164b34c75fc75c4_img.jpg\) Share](#)

May 2022: This post was reviewed and updated with more details like using EXPLAIN ANALYZE, updated compression, ORDER BY and JOIN tips, using partition indexing, updated stats (with performance improvements), added bonus tips.

[Amazon Athena](#) is an interactive query service that makes it easy to analyze data stored in [Amazon Simple Storage Service](#) (Amazon S3) using standard SQL. Athena is serverless, so there is no infrastructure to manage, and you pay only for the queries that you run. Athena is easy to use. Simply point to your data in Amazon S3, define the schema, and start querying using standard SQL.

<https://aws.amazon.com/blogs/big-data/top-10-performance-tuning-tips-for-amazon-athena/>

3.4 Workgroups

Athena Workgroups

Athena Workgroups are used to isolate queries between different teams, workloads or applications, and to set limits on amount of data each query or the entire workgroup can process

Workload Isolation

Query Metrics

Cost Controls

Workgroups – Workload Isolation

Unique query output location per Workgroup

Encrypt results with unique AWS KMS key per Workgroup

Workgroup name* Use 1 - 128 characters. (A-Z,a-z,0-9,_,-,.)

Description Use up to 1024 characters.

Query result location Select Enter a path to an S3 bucket or prefix.

Encrypt query results Encrypt results stored in S3

Encryption type i

Encryption key i Create KMS key

Metrics Publish query metrics to AWS CloudWatch i

Override user settings i

Collect and publish aggregated metrics per Workgroup to AWS CloudWatch

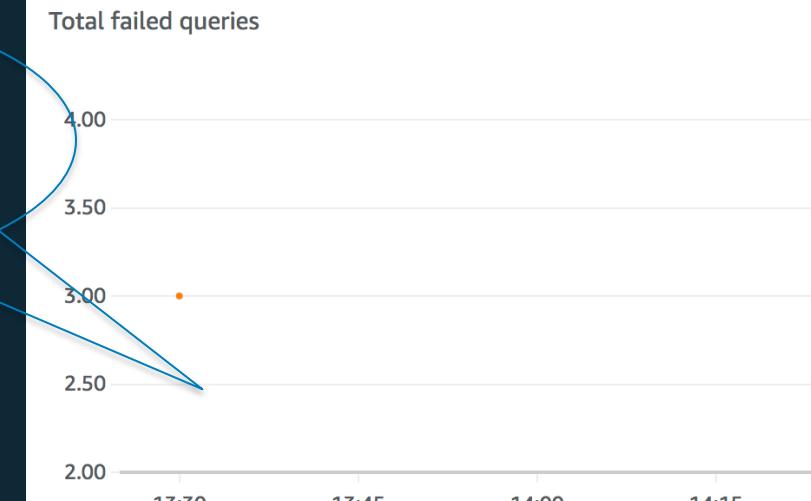
Use Workgroup settings eliminating need to configure individual users

Workgroups – Metric Reporting

Total bytes scanned per Workgroup



Total failed queries per Workgroup

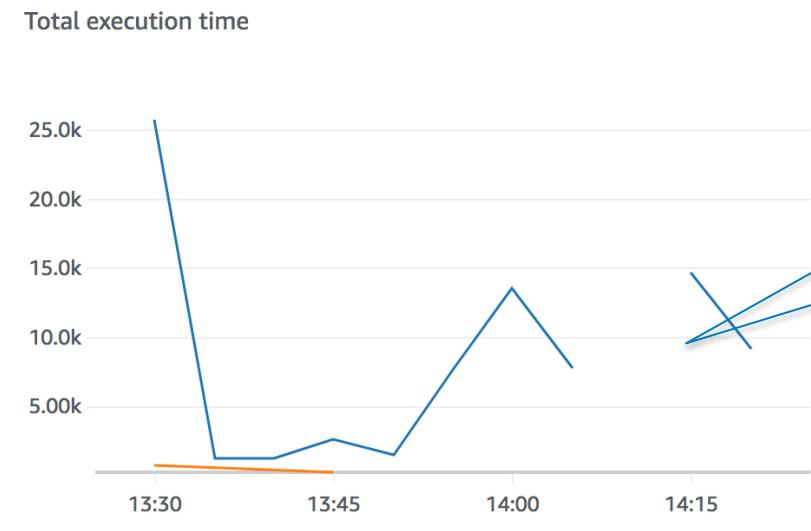


Total succeeded queries



Total successful queries per Workgroup

Total execution time



Total query execution time per Workgroup

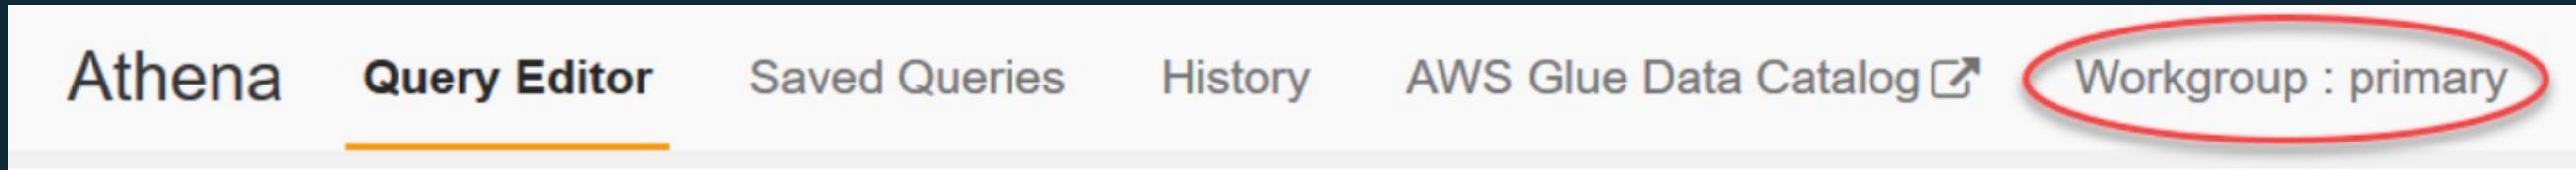
Workgroups – Cost Controls

- Per query data scanned threshold; exceeding, will cancel query
- Trigger alarms to notify of increasing usage and cost
- Disable Workgroup when all queries exceed a maximum threshold

Any Athena metric: successful/failed & total queries, query run time, etc.

Data limit	Time period	Action	
10 Gigabytes	Not applicable	Query will be cancelled.	
1 Terabytes	24 hours	Send notification to topic : arn:aws:sns:us-east-1:9	9:AthenaAlarm
10 Gigabytes	1 hour	Send notification to topic : arn:aws:sns:us-east-1:9	9:AthenaAlarm

Default Workgroup



By default, if you have not created any workgroups, all queries in your account run in the primary workgroup

3.5 Views

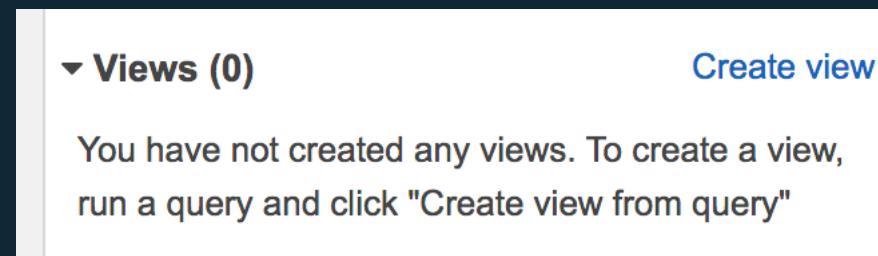
View

- A view in Amazon Athena is a logical, not a physical table.
- The query that defines a view runs each time the view is referenced in a query.
- Views do not contain any data and do not write data.
 - Instead, the query specified by the view runs each time you reference the view by another query.

Create a View

- Creates a new view from a specified SELECT query.
 - The view is a logical table that can be referenced by future queries.
- Flow:

- Create View:



- Write View:

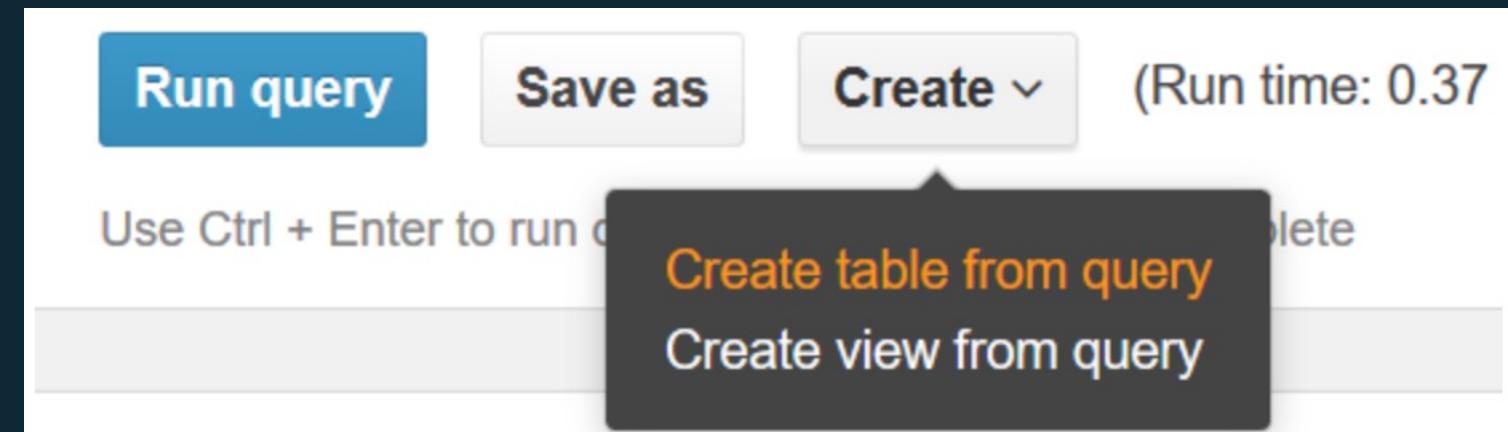
```
1 -- View Example
2 CREATE VIEW test AS
3 SELECT
4 orderkey,
5 orderstatus,
6 totalprice / 2 AS half
7 FROM orders
```

The screenshot shows a query editor with two tabs: 'New query 1' and 'New query 2'. The second tab is active and contains the following SQL code:
1 -- View Example
2 CREATE VIEW test AS
3 SELECT
4 orderkey,
5 orderstatus,
6 totalprice / 2 AS half
7 FROM orders

3.6 CTAS

CREATE TABLE AS SELECT (CTAS)

- Creates a new table in Athena from the results of a SELECT statement from another query
- Athena stores data files created by the CTAS statement in a specified location in Amazon S3



Use CTAS queries to:

- Create tables from query results in one step, without repeatedly querying raw data sets. This makes it easier to work with raw data sets.
- Transform query results into other storage formats, such as Parquet and ORC. This improves query performance and reduces query costs in Athena.
- Create copies of existing tables that contain only the data you need.

Extract, Transform and Load data into S3 data lake using CTAS and INSERT INTO statements in Amazon Athena

by Pathik Shah | on 26 NOV 2019 | in [Amazon Athena](#), [Analytics](#), [AWS Big Data](#) | [Permalink](#) | [Comments](#) | [Share](#)

[Amazon Athena](#) is an interactive query service that makes it easy to analyze the data stored in [Amazon S3](#) using standard SQL. Athena is serverless, so there is no infrastructure to manage, and you pay only for the queries that you run. You can reduce your per-query costs and get better performance by compressing, partitioning, and converting your data into columnar formats. To learn more about best practices to boost query performance and reduce costs, see [Top 10 Performance Tuning Tips for Amazon Athena](#).

Overview

This blog post discusses how to use Athena for extract, transform and load (ETL) jobs for data processing. This example optimizes the dataset for analytics by partitioning it and converting it to a columnar data format using Create Table as Select (CTAS) and INSERT INTO statements.

[CTAS](#) statements create new tables using standard SELECT queries to filter data as required. You can also partition the data, specify compression, and convert the data into columnar formats like Apache Parquet and Apache ORC using CTAS statements. As part of the execution, the resultant tables and partitions are added to the [AWS Glue](#) Data Catalog, making them immediately available for subsequent queries.

[INSERT INTO](#) statements insert new rows into a destination table based on a SELECT query statement that runs on a source table. If the source table's underlying data is in CSV format and destination table's data is in Parquet format, then INSERT INTO can easily transform and load data into destination table's format. CTAS and INSERT INTO statements can be used together to perform an initial batch conversion of data as well as incremental updates to the existing table.

<https://aws.amazon.com/blogs/big-data/extract-transform-and-load-data-into-s3-data-lake-using-ctas-and-insert-into-statements-in-amazon-athena/>

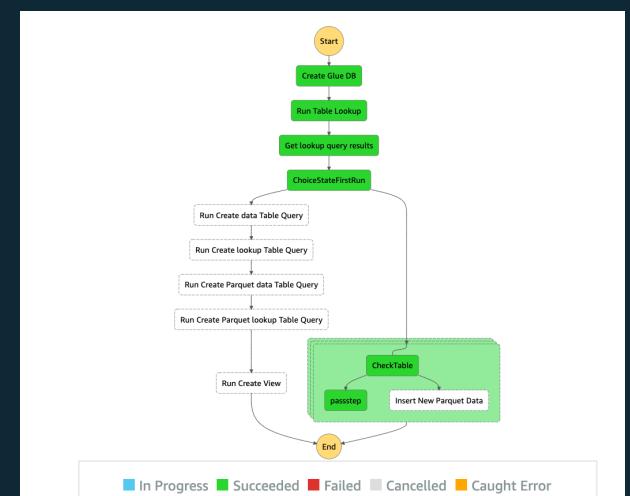
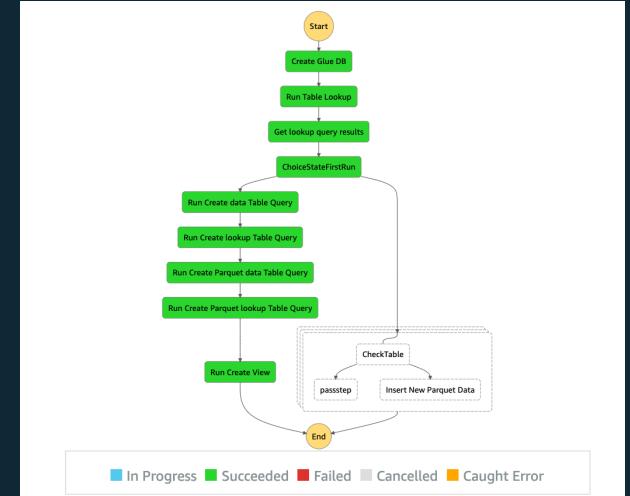
Build and orchestrate ETL pipelines using Amazon Athena and AWS Step Functions

by Behram Irani, Dipankar Kushari, and Rahul Sonawane | on 12 OCT 2021 | in [Amazon Athena](#), [AWS Step Functions](#) | [Permalink](#) | [Comments](#) | [Share](#)

Extract, transform, and load (ETL) is the process of reading source data, applying transformation rules to this data, and loading it into the target structures. ETL is performed for various reasons. Sometimes ETL helps align source data to target data structures, whereas other times ETL is done to derive business value by cleansing, standardizing, combining, aggregating, and enriching datasets. You can perform ETL in multiple ways; the most popular choices being:

- Programmatic ETL using Apache Spark. [Amazon EMR](#) and [AWS Glue](#) both support this model.
- SQL ETL using Apache Hive or PrestoDB/Trino. Amazon EMR supports both these tools.
- Third-party ETL products.

Many organizations prefer the SQL ETL option because they already have developers who understand and write SQL queries. However, these developers want to focus on writing queries and not worry about setting up and managing the underlying infrastructure.



<https://aws.amazon.com/blogs/big-data/build-and-orchestrate-etl-pipelines-using-amazon-athena-and-aws-step-functions/>

3.7 Monitoring & Auditing

Workgroups

- Separate users, teams, applications, or workloads
- Enable queries on Requester Pays buckets in Amazon S3
- In order to:
 - Set limits on amount of data each query or the entire workgroup can process
 - Track costs
 - Decide which workgroups to create.
 - Create workgroups as needed and add tags to them.
 - Create IAM policies for your users, groups, or roles to enable their access to workgroups.
 - Specify a location in Amazon S3 for query results and encryption settings.

Monitoring Athena Queries with CloudWatch Metrics

- Use CloudWatch Events with Athena:
 - EngineExecutionTime – in milliseconds
 - ProcessedBytes – the total amount of data scanned per DML query
 - QueryPlanningTime – in milliseconds
 - QueryQueueTime – in milliseconds
 - ServiceProcessingTime – in milliseconds
 - TotalExecutionTime – in milliseconds, for DDL and DML queries

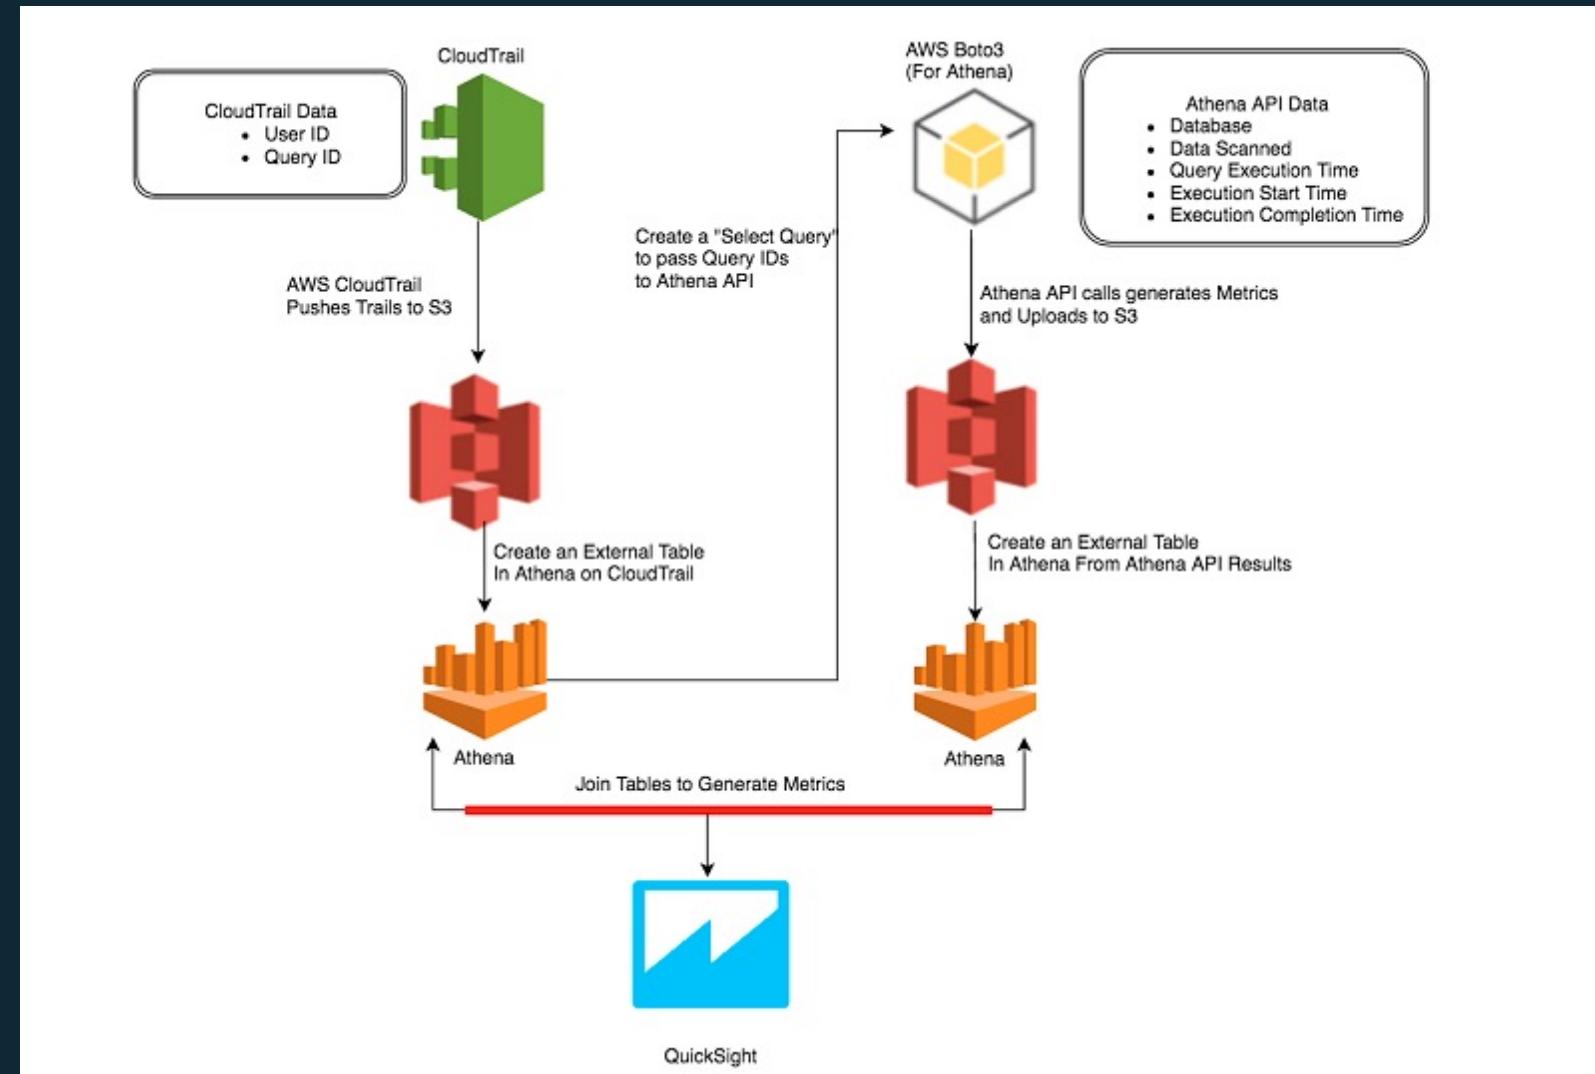
Monitoring Athena Queries with CloudWatch Events

- Use simple rules to match events and route them to one or more target functions or streams.
- Respond to operational changes and takes corrective action as necessary, by
 - Sending messages to respond to the environment
 - Activating functions
 - Making changes
 - Capturing state information

```
{  
  "source": [  
    "aws.athena"  
  ],  
  "detail-type": [  
    "Athena Query State Change"  
  ],  
  "detail": {  
    "currentState": [  
      "SUCCEEDED"  
    ]  
  }  
}
```

CloudTrail

Monitor Athena with AWS CloudTrail



4. Summary

Key Benefits of Athena

- Fast ad-hoc queries
- Query directly against data in S3
- Use standard ANSI SQL
- No infrastructure to setup and manage
- Use IAM for security
- JDBC driver for BI tools and SQL clients
- Pay per query, depending on data scanned



Thank you!

Michael Lin

linmicht@amazon.com